

Vivek Karuppuswamy and the DeepHow Hiring Team
Re: Product Director, Physical AI Platform

DeepHow is at an unusually important product inflection point. Knowledge Capture & Transfer, Operational Output Verification, Live SOP Verification, and Time & Motion AI can create substantially more value together than as neighboring capabilities—but only if the platform makes the performed unit of work legible across all four.

That is the work I would be excited to lead: defining what must be shared, what should remain federated, and how DeepHow compresses the path from first access to first verified value. I made that reasoning visible at russelldudek.github.io/DeepHow/. The candidate vision includes a Work Address Resolver, a 120-day entry plan, and a Work Address Review that connect expert knowledge, expected action, observed execution, evidence thresholds, and learning return. It is not a claim about undisclosed DeepHow process; it is a transparent sample of how I would frame the mandate and make product decisions.

My background is unusually grounded in both the physical work and the mechanisms around it. At Vape-Jet, I lead AI-first operating transformation across production, quality, support, engineering issue flow, enterprise systems, telemetry, and a field fleet of approximately 1,000 software-enabled robots. Earlier, I led high-reliability engineering and manufacturing at Compunetics, customer-backward launch mechanisms at Amazon, regulated autonomous drone operations at ZeusVu, and direct AI-readiness and workflow design through DudeWorth.

My path is not a conventional enterprise-software product ladder; it is grounded in the boundary DeepHow is productizing. I have repeatedly worked where software, machines, quality systems, operators, customers, and commercial commitments meet—and where product decisions succeed only when they survive real operating conditions.

That combination shapes my product judgment. An AI capability does not earn scale because the demo is compelling. It earns scale when the customer outcome is explicit; data and work context are usable; confidence, latency, human authority, and recovery are calibrated; builders can understand the platform; and in-market evidence tells the team what to standardize, change, or stop.

In the first 120 days, I would listen before prescribing, map the product lines to a shared work object and evidence vocabulary, identify where platform leverage is already emerging, and select a bounded cross-product proof with a real customer outcome. I would then install the decision cadence, evaluation posture, ownership model, and scorecard needed to turn the proof into a repeatable platform capability rather than another isolated success.

I would bring a practical operating cadence to DeepHow: listen deeply to frontline users, customers, partners, GTM, design, and engineering; write crisp product decisions; run bounded experiments; instrument trust and value; and convert successful deployments into reusable platform contracts. I can speak with the CIO, AI team, quality leader, industrial engineer, and operator without flattening their different truths into one generic story.

DeepHow's public work with Foxconn and NVIDIA shows what becomes possible when approved procedures, real-time observation, AI reasoning, and operational outcomes close the loop. I would welcome the opportunity to help make that loop a coherent, builder-friendly, enterprise-grade platform.

Sincerely,
Russell Dudek